

ENTRE
PRENEURS
FIRST



n NOZOMIO HACKATHON 9th May

▲ Vercel ☁ Codex ⚙️ Devin

◀ InsForge 🔌 Tensorlake 📄

convex ⚡ Hyperspell ⚡ Aside

Document outline

Document outline.....	2
Discord.....	3
Schedule.....	4
About the sponsors.....	4
Prizes.....	7
Track specific prizes.....	7
Hyperspell track.....	7
Event miscellaneous.....	8
Wifi.....	8
On the day Credits.....	8
Tracks.....	10
📡 Always-On Agents 📡.....	10
🚀 Ship It - Full-Stack Agents 🚀.....	12
🎯 AI-Native Growth Tools 🎯.....	14
🧠 The Company Brain 🧠.....	16
Submission.....	18
Judging Format.....	18
Judging Method.....	18
Rules.....	19

Discord

For all event information, team search and updates please join our discord:

<https://discord.gg/KSERFyh6>

Schedule

8am: Doors open & Breakfast is served

8:30am: Nozomio & Sponsors introductory talk

9:15am: Hacking starts

12:00am: Hyperspell speaker session

12:30pm: Lunch is served

3:00pm: Speaker session

6:00pm: Submissions close & dinner is served

6:10pm: In-person judging starts

7:30pm: Judging concludes & Winners announced

8:00pm: Relax

About the sponsors

About [Nozomio](#)

Nozomio is a product and research lab solving context in the age of AI. Its first product, Nia, is a search and indexing API that continuously provides context from docs, research papers, datasets, codebases, and more, so agents never rely on stale data. It is scalable, 5x cheaper, and reliable.

About [AI Nexus](#)

AI Nexus is an event agency with a global community of 5,000+ builders passionate about AI. They run hackathons and tech events across 6 locations worldwide, currently, with the leading AI innovators.

About [Vercel](#)

Vercel is the platform for frontend developers, providing the speed and reliability innovators need to create at the moment of inspiration. Vercel enables teams to ship better websites, faster and with zero configuration deployments, global edge infrastructure, and seamless Git integration.

About [InsForge](#)

InsForge is the backend built for agentic development. It gives AI coding agents everything they need to ship fullstack apps such as Postgres databases, authentication, storage, edge functions, and AI model access - with no setup or server management required. Proven to make agents 1.7x more accurate and 30% more token-efficient.

About [Reacher](#)

Reacher is an AI-powered creator marketing platform built for social commerce. It automates the full affiliate workflow - from creator discovery and outreach to campaign management and analytics - making marketing teams 10x more productive on TikTok Shop, Instagram, and YouTube. Backed by Y Combinator (S25).

About [Hyperspell](#)

Hyperspell is the company brain for AI agents. It connects to a team's data (Slack, email, docs, code, CRM) and synthesizes it into context any agent can read. Backed by Y Combinator and SOC 2 certified.

About [Tensorlake](#)

Tensorlake is a stateful sandbox compute platform built for AI agents. It provides fast, isolated execution environments where agents and tool calls can run with their own filesystem, shell, and software stack. With 150ms startup times and industry-leading file system performance, Tensorlake is built for complex, long-running agentic workloads at scale.

About [Convex](#)

Convex is the backend platform that keeps your app in sync. It lets developers express their entire backend - databases, queries, auth, and APIs - in pure TypeScript, with real-time updates that just work out of the box. No state managers, no cache invalidation, no infrastructure headaches.

About [Aside](#)

Aside is building a new kind of browser, one that acts as the operating system for AI. Instead of forcing agents to rely on endless integrations and constant user approval, Aside lets AI work directly inside the browser where modern work already happens: docs, banking, recruiting, design, and more. Backed by Y Combinator and based in San Francisco.

About [Devin](#)

Devin is an AI coding agent and software engineer that helps developers build better software faster. Parallel cloud agents for serious engineering teams.

About [OpenAI Codex](#) by TatianaSF (Codex Ambassador)

OpenAI's coding agent that helps teams build and ship faster. OpenAI Codex is a powerful AI-driven software engineering assistant that can read, generate, edit, and even run code. Learn more at <https://openai.com/codex/>.

Venue provided by [EF](#)

Entrepreneurs First (EF) turns exceptional individuals into funded founders. Some might have an early idea, others may need a cofounder. What matters is their potential. We back these founders first, then give them a bridge to Silicon Valley. To date, their companies are worth > \$16B.

Prizes

1st Place:

- M5 Macbook Pros
- Guaranteed job interview with Arlan
- Various Credits

2nd Place:

- M4 Mac Mini
- Various Credits
- \$500

3rd Place:

- AirPods Pros
- Various Credits
- \$300

Top 10:

1 Month of Hinge premium

Other prizes:

- First Class trip to Las Vegas: Gamble and raise your pre-seed
- Job interviews with the various sponsors
- Stussy gear

Track specific prizes

Hyperspell track

Prize for best submission on this track:

- \$1k cash
- 6 months free Hyperspell unlimited usage
- Working session with the Hyperspell founders (Conor + Manu) to deploy it
- Amplification on Hyperspell channels (Twitter, LinkedIn newsletter)

Event miscellaneous

Wifi

Guest Wifi: Entrepreneur Firsts-Guest

Guest Password: 501Folsom!

On the day Credits

Nia:

TBD

InsForge

\$100 credits for each user

Step 1: Visit <https://insforge.dev/promo/NIA> and claim your pro plan

Tensorlake

Free access for the entire day on May 9th

Reacher

Free access for the entire day on May 9th

Step 1: Visit  Reacher Nozomio Hackathon Setup.pdf

Convex

Free deployments - no credit card required - more info [here](#)

Hyperspell

Free access for the entire day on May 9th

Vercel v0

\$30 v0 credits for the first 200 people

Step 1: Use code NOZOMIO-V0 on the v0 platform to redeem

OpenAI Codex & OpenAI API

\$50 credits towards each platform

Speak to Tatiana to retrieve your Codex credits.

Devin

\$100 Devin coupon to all attendees!

For Devin credits , please fill out the form:

<https://forms.gle/USxbQJxoPQ1AGMG4A>

Step 1: Sign up with your registration email - <https://app.devin.ai/signup>

Step 2: Create your Devin Account so that the credits can be added

Step 3: Instructions on how to gain your credits will be sent via email once you have completed your form

Tracks

Always-On Agents

Sponsor: Nia + Tensorlake

Build an agent that runs continuously in the background, remembers what it has learned across sessions, and acts without a human prompting it. The bar is high: removing either the background execution or the stateful memory should break your demo.

The dream shape is an agent that wakes up on its own and knows everything it has done before - a research monitor that learns your reading history over weeks, a support agent with durable per-customer memory across months, or a dev assistant that reviews PRs overnight and has internalised your codebase's conventions.

Examples of what to build

- A research agent that monitors a topic continuously, tracks what you've read and dismissed, and sends weekly briefs that get sharper every week
- A support agent that handles tickets 24/7 with durable per-customer memory
- An overnight PR reviewer that clears your team's queue before standup
- A personal agent that runs on your calendar and email, remembers what you asked last Tuesday, and learns your preferences over time

Sponsors

1. 1 Nia is a context augmentation layer and MCP server that indexes repos, documentation sites, PDFs, Slack, Google Drive, local files - anything - so agents always have up-to-date, accurate context rather than hallucinating from stale training data. In this track, Nia keeps the agent's knowledge fresh and current across every run.

What they provide:

- MCP server with 15+ specialised tools for semantic search across indexed sources
 - Indexes entire codebases, docs sites, research papers, datasets, and local folders
 - Nia Sync daemon for continuously syncing local data (files, iMessage, databases)
 - Nia Oracle - autonomous research agent for multi-source research workflows
 - Cross-session context: save findings in one agent, retrieve in another
 - Proven to reduce hallucination rates by 11.3% vs leading alternatives
2. Tensorlake provides the infrastructure for background and stateful agent execution. Agents built with Tensorlake run on their own - triggered by schedules, webhooks, or other agents - and remember what they have done across runs. Durable memory is a first-class primitive, not bolted on as a RAG afterthought.

What they provide:

- Background execution: agents triggered by schedules, webhooks, events, or other agents
- Stateful execution: durable memory that survives invocations, sessions, and restarts
- Sandbox environment for real work - code, APIs, data transforms
- Full agentic infrastructure: the agent runs while no one is watching

Judging criteria

Criterion	1 - Poor	2 - Weak	3 - Solid	4 - Strong	5 - Exceptional
Genuine Background Execution 30%	Triggered manually or by user input only	Automated trigger exists but relies on human setup each time	Runs on a schedule or event without user involvement	Fully autonomous triggers; handles failures and retries without human intervention	Deeply autonomous; multi-source triggers, graceful recovery, runs reliably over hours
Statefulness 25%	No memory between runs; starts fresh every time	Saves some output but doesn't use it to change future behaviour	Prior state influences current run in a meaningful way	Rich durable memory; agent demonstrably improves or adapts based on history	Memory is load-bearing; removing it would break the demo entirely
Agentic Depth 25%	Single-step, no decision-making	Basic branching but no real planning or recovery	Plans multi-step flows; handles some edge cases	Retries on failure; adapts plan mid-run; reasons about what to do next	Full agentic loop: plans, executes, reflects, recovers, and improves autonomously
Demo & Presentation 10%	Hard to follow what the agent does	Shows it works but misses the compelling case	Clear demo; explains the value well	Live demo lands; judges want to try it themselves	Story + demo make the case unforgettably
Judge's Personal Rating 10%	Not interesting or compelling	Somewhat interesting but forgettable	Solid idea; I can see the appeal	I'd tell a colleague about this	This genuinely excites me; I want to see it succeed

Ship It - Full-Stack Agents

Sponsors: Nia + InsForge

Build a complete, production-deployed agentic application with real authentication, a live backend, and AI logic that does not break when users go off-script. The challenge: get it actually shipped and running, not just demo-ready.

Not localhost. Not a polished UI with no backend. Something a stranger can open, sign up for, and use today. Hyperspell's context graphs handle agent intelligence, InsForge gets it shipped end-to-end, and Nia ensures the agent never hallucinates on your own codebase or docs.

Examples of what to build

- A deployable coding assistant that understands your private codebase via Nia and ships features autonomously
- A fully live SaaS tool with auth, database, and real user flows - powered by an agent that handles the messy edge cases
- A multi-agent app where different agents (planning, execution, review) share context via Nia and Hyperspell
- Anything production-grade: if someone outside your team can use it end-to-end, you're on the right track

Sponsors

1. Nia is a context augmentation layer and MCP server that indexes repos, documentation sites, PDFs, Slack, Google Drive, local files - anything - so agents always have up-to-date, accurate context rather than hallucinating from stale training data. In this track, Nia keeps the agent's knowledge fresh and current across every run.

What they provide:

- MCP server with 15+ specialised tools for semantic search across indexed sources
 - Indexes entire codebases, docs sites, research papers, datasets, and local folders
 - Nia Sync daemon for continuously syncing local data (files, iMessage, databases)
 - Nia Oracle - autonomous research agent for multi-source research workflows
 - Cross-session context: save findings in one agent, retrieve in another
 - Proven to reduce hallucination rates by 11.3% vs leading alternatives
2. InsForge is about production, not demos. Set up authentication, database, and deployment - not just UI, not just localhost. Your app should be usable end-to-end by someone outside your team by the end of the day.

What they provide:

- Production-ready backend scaffolding with auth and database out of the box
- Live deployment so your app is accessible to real users on the day
- Focus on reliability and real user flows, not just a polished front end

3. Hyperspell provides context graphs for building sophisticated AI agents. Use Hyperspell to wire up agents like Claude Code, Hermes, or OpenClaw with structured, persistent context about your project - so agents understand intent and architecture, not just individual prompts.

What they provide:

- Context graphs that give agents structured understanding of your project
- Supports leading agent frameworks: Claude Code, Hermes, OpenClaw, and more
- Enables multi-agent workflows where context flows between agents coherently

Judging criteria

Criterion	1 - Poor	2 - Weak	3 - Solid	4 - Strong	5 - Exceptional
Production Readiness 35%	localhost only; nothing deployed	Deployed but missing auth or database; not usable by others	Live URL, auth, and database all working; usable by someone outside the team	Polished and reliable; handles edge cases and real user flows without breaking	Could be handed to a real user today; impressively complete for 24 hours
Agent Reliability 30%	Agent breaks on anything outside the happy path	Works on happy path; falls over on edge cases	Handles a few off-script scenarios; recovers from some errors	Robust across multiple flows; fails gracefully and explains errors well	Agent holds up under adversarial or unexpected input; feels production-grade
Full-Stack Depth 25%	UI only; no real backend or AI logic	Frontend + some backend, but AI is a thin wrapper	Frontend, backend, and AI logic all present and connected	All layers well-integrated; AI drives real functionality, not just surface features	Seamless full stack; each layer non-trivially implemented and AI is core
Demo & Presentation 10%	Hard to follow what the product does	Shows it works but misses the compelling case	Clear demo; explains the value well	Live demo lands; judges want to try it themselves	Story + demo make the case unforgettably
Judge's Personal Rating 10%	Not interesting or compelling	Somewhat interesting but forgettable	Solid idea; I can see the appeal	I'd tell a colleague about this	This genuinely excites me; I want to see it succeed

AI-Native Growth Tools

Sponsors: Nia + Reacher

Build tools that do revenue work autonomously - creator scouting, outreach automation, campaign building, or attribution analysis. Think beyond dashboards: the best submissions will have the agent doing the work, not just reporting on it.

Reacher provides per-team TikTok Shop demo data, market-wide Social Intelligence, and sandboxed write endpoints so agents can create real campaigns, sample requests, and outreach drafts without any external dispatch. Nia adds the context layer - index market reports, creator briefs, or campaign history so your agent reasons with the full picture, not just the latest query.

Examples of what to build

- A TikTok creator scouting agent: brief → ranked list with evidence, searching beyond a single network
- An always-on market monitor: continuous watch on category and trending-video shifts, with alerts on competitor creator spikes
- An end-to-end TikTok Shops campaign builder: discovery → vetting → outreach creation → follow-up, fully sandboxed
- Attribution forensics: explain week-over-week GMV changes using creator, video, and category signals

Sponsors

1. Nia is a context augmentation layer and MCP server that indexes repos, documentation sites, PDFs, Slack, Google Drive, local files - anything - so agents always have up-to-date, accurate context rather than hallucinating from stale training data. In this track, Nia keeps the agent's knowledge fresh and current across every run.

What they provide:

- MCP server with 15+ specialised tools for semantic search across indexed sources
 - Indexes entire codebases, docs sites, research papers, datasets, and local folders
 - Nia Sync daemon for continuously syncing local data (files, iMessage, databases)
 - Nia Oracle - autonomous research agent for multi-source research workflows
 - Cross-session context: save findings in one agent, retrieve in another
 - Proven to reduce hallucination rates by 11.3% vs leading alternatives
2. Reacher provides the TikTok Shop data infrastructure for this track. Every team gets a sandboxed demo environment with realistic data - no payment card, no onboarding friction.

What they provide:

- MCP server at api.reacherapp.com/mcp with 33 tools covering creators, products, videos, samples, GMV metrics, and the full Social Intelligence market catalogue
- Social Intelligence: market-wide creator, seller, and trending-video data across the entire TikTok Shop ecosystem

- Per-team demo dataset - a scoped clone of a dummy account with creators, videos, GMV history, and samples. Realistic enough for multi-step agent reasoning, isolated so teams don't interfere
- Sandboxed write endpoints - POST /automations, /samples/request, /outreach/draft persist records for demo purposes but never dispatch to TikTok or trigger real emails
- One-click MCP connector in Claude.ai, Cursor, and any MCP client

Judging criteria

Criterion	1 - Poor	2 - Weak	3 - Solid	4 - Strong	5 - Exceptional
Depth of Social Intelligence Usage 30%	Doesn't use Reacher data meaningfully; surface-level or none	Uses one or two tools but misses the richness of the data available	Draws on multiple data types (creators, GMV, videos) in a coherent way	Social Intelligence is central to how the agent reasons and makes decisions	Couldn't exist without Reacher; extracts genuine insight a human would miss
Agentic Complexity 30%	Single API call wrapped in a UI; no real agent logic	Some chaining but no real decision-making or recovery	Multi-step flow with basic branching and error handling	Plans, retries, and adapts; handles off-script prompts without breaking	Full agentic loop with autonomous decision-making across the entire flow
End-to-End Flow 20%	Produces a chart or report; agent does no real work	Agent does some work but flow is incomplete or requires manual steps	Full flow works end-to-end in the sandbox; campaigns or outreach drafts created	Seamless flow from discovery to output; sandboxed writes demonstrate real capability	Agent ships a complete campaign autonomously; demo feels like a real product
Demo & Presentation 10%	Hard to follow what the agent does	Shows it works but misses the compelling case	Clear demo; explains the value well	Live demo lands; judges want to try it themselves	Story + demo make the case unforgettably
Judge's Personal Rating 10%	Not interesting or compelling	Somewhat interesting but forgettable	Solid idea; I can see the appeal	I'd tell a colleague about this	This genuinely excites me; I want to see it succeed

The Company Brain

Sponsors: Nia + Hyperspell

Build an agent that synthesizes context across a company's data and does something genuinely valuable with it. The best submissions won't just answer questions, they'll do work. Sales agents that know every deal in the pipeline. Support agents that know every past ticket. Marketing agents that have internalized the brand voice. Coding agents that know the codebase, the team's conventions, and last week's design review.

Hyperspell gives you ingestion (Slack, Gmail, Drive, GitHub, Notion, and more) and search across all of it. You decide what the brain looks like and what agents do with the context.

Examples of what to build

- A sales agent that drafts deal-specific outreach by synthesizing CRM, email threads, and past won deals
- A marketing agent that ships campaigns matching brand voice learned from every doc and Slack the team has ever written
- An onboarding agent that knows the company's tech stack, Slack history, and processes better than any human
- A support agent that resolves tickets with full context across product decisions, past tickets, and customer history

Sponsors

1. Nia is a context augmentation layer and MCP server that indexes repos, documentation sites, PDFs, Slack, Google Drive, local files - anything - so agents always have up-to-date, accurate context rather than hallucinating from stale training data. In this track, Nia keeps the agent's knowledge fresh and current across every run.

What they provide:

- MCP server with 15+ specialised tools for semantic search across indexed sources
 - Indexes entire codebases, docs sites, research papers, datasets, and local folders
 - Nia Sync daemon for continuously syncing local data (files, iMessage, databases)
 - Nia Oracle - autonomous research agent for multi-source research workflows
 - Cross-session context: save findings in one agent, retrieve in another
 - Proven to reduce hallucination rates by 11.3% vs leading alternative
2. Hyperspell is the company brain for AI agents. They ingest all of an organization's data and synthesizes it into context any agent can read.

What they provide:

- Ingestion across Slack, Gmail, Google Drive, GitHub, Notion, and more
- Semantic search across all sources, synthesis, and presentation as a filesystem for all agents

- Built for agents like Claude Code, Hermes, and OpenClaw

Judging criteria

Criterion	1 - Poor	2 - Weak	3 - Solid	4 - Strong	5 - Exceptional
Cross-Source Synthesis 30%	Single source or surface-level use; no real synthesis	Two sources but stitched, not synthesised	Combines 3+ sources to produce coherent context	Synthesis is central; agent reasons across sources fluidly	Produces context no single source could; the brain is the product
Real Work, Not Just Answers 25%	Q&A bot or chatbot wrapper	Surfaces info but a human still does the work	Agent does part of the work end-to-end	Agent ships work a human would actually use	Agent delivers output that replaces hours of human work
Hyperspell Integration Depth 25%	Hyperspell used as a thin wrapper or not at all	Surface-level use of ingestion or search	Hyperspell is meaningfully load-bearing	Synthesis through Hyperspell is core to how the agent works	Removing Hyperspell would break the demo entirely
Demo & Presentation 10%	Hard to follow what the agent does	Shows it works but misses the compelling case	Clear demo; explains the value well	Live demo lands; judges want to try it themselves	Story + demo make the case unforgettably
Judge's Personal Rating 10%	Not interesting or compelling	Somewhat interesting but forgettable	Solid idea; I can see the appeal	I'd tell a colleague about this	This genuinely excites me; I want to see it succeed

Submission

Please read before submitting:

- **One submission per team.** Only one team member should fill out this form. Duplicate submissions will be disqualified.
- **Submissions close at 6:00pm sharp.** Late submissions may not be accepted.
- **You must be present for judging.** At least one team member must be available for the in-person judging session starting at 6:10pm.
- **Have these ready before you start:** your deployed demo link, GitHub repo URL, and the names and emails of all team members.
- **Incomplete submissions may be disqualified.** Make sure every field is filled out - especially the demo link. A localhost link is not a valid demo.

Judging Format

Step 1: Submit your project here <https://forms.gle/fkoEXRo3L2MVkkz87>

Step 2: Judging will be conducted **in person** first. Each team will receive 3 mins to present their project directly to the judges. Judges will individually score and the average will be marked as the final grade.

Step 3: The top 6 across all tracks will present live in front of the entire room and in front of the final judges.

Judging Method

There are 4 tracks. Each track is judged separately by the judges. However, the final rankings are track-agnostic. The top 6 finalists and top 3 winners are determined purely by score, regardless of which track you submitted to. This means you're ultimately competing against everyone, not just teams in your track.

We understand this may seem unusual but we've designed it this way to remove bias and ensure the best projects win, period.

Judges score individually and the average is your final grade.

Rules

1. Max of 3 participants per team, solo is also welcomed
2. No re-using personal projects
3. Project must be submitted before 6pm
4. Any form of cheating or rule-breaking disqualifies the team immediately
5. Organisers reserve the right to disqualify any project that violates the spirit of the event
6. Winners must be present at the awards to claim prizes
7. Prize fulfilment handled by sponsors directly - timelines may vary
8. Teams can only join and present for 1 track